

Dynamic Programming*

Xiaofeng Gao

Department of Computer Science and Engineering
Shanghai Jiao Tong University, P.R.China

Algorithm Course: Shanghai Jiao Tong University

* Special thanks is given to Prof. Kevin Wayne@Princeton for sharing his slides, and also given to Mr. Chao Wang from CS2014@SJTU and Mr. Hongjian Cao from CS2015@SJTU for producing this lecture.

Algorithmic Paradigms

Greedy: Build up a solution incrementally, myopically optimizing some local criterion.

Divide-and-conquer: Break up a problem into sub-problems, solve each sub-problem independently, and combine solution to sub-problems to form solution to original problem.

Dynamic programming: Break up a problem into a series of overlapping sub-problems, and build up solutions to larger and larger sub-problems.

Outline

1 Introduction

- Background
- Introductory Example: Weighted Interval Scheduling

2 Popular Recipes

- Segmented Least Squares
- Knapsack Problem
- RNA Secondary Structure

3 Hirschberg's Alignment Algorithm

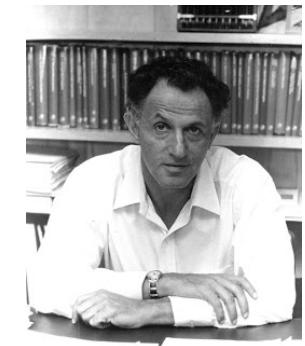
- String Similarity
- Sequence Alignment in Linear Space

History

Richard E. Bellman (1920-1984): Pioneered the systematic study of dynamic programming in 1950s.

Etymology:

- Dynamic programming = planning over time
- Secretary of Defense had pathological fear of mathematical research.
- Bellman sought a “dynamic” adjective to avoid conflict.



Applications

Areas: Bioinformatics, Control Theory, Information Theory, Operations Research, Computer Science (Theory, Graphics, AI, Compilers, Systems, ···)

Some Famous Algorithms

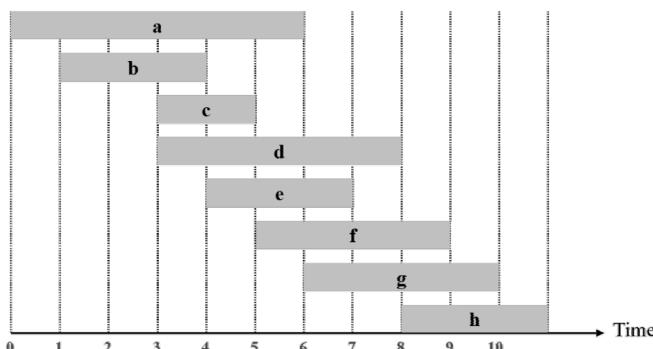
- Avidan-Shamir for seam carving.
- Unix diff for comparing two files.
- Viterbi for hidden Markov models.
- De Boor for evaluating spline curves.
- Knuth-Plass for word wrapping text in T_EX.
- Smith-Waterman for genetic sequence alignment.
- Bellman-Ford-Moore for shortest path routing in networks.
- Cocke-Kasami-Younger for parsing context free grammars.
- Needleman-Wunsch/Smith-Waterman for sequence alignment.

Weighted Interval Scheduling Problem

Job j starts at s_j , finishes at f_j , and has weight or value $w_j > 0$.

Two jobs are **compatible** if they don't overlap.

Goal: find maximum weight subset of mutually compatible jobs.



Dynamic Programming Books

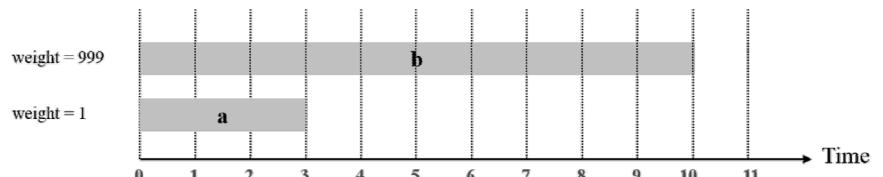


Unweighted Interval Scheduling Review

Recall: Greedy algorithm works if all weights are 1.

- Consider jobs in ascending order of finish time.
- Add job to subset if it is compatible with previously chosen jobs.

Observation: Greedy algorithm can fail spectacularly if arbitrary weights are allowed.

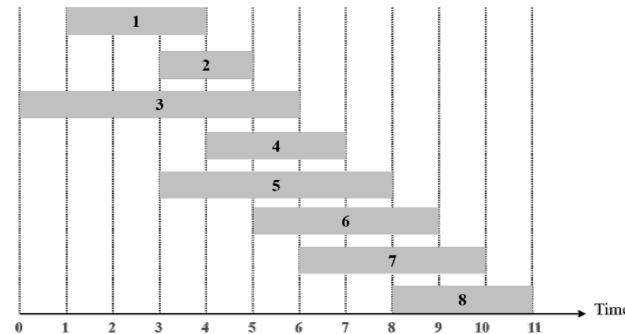


Weighted Interval Scheduling

Notation: Label jobs by finishing time: $f_1 \leq f_2 \leq \dots \leq f_n$.

Definition: $p(j)$ = largest index $i < j$ such that job i is compatible with j .

Example: $p(8) = 5, p(7) = 3, p(2) = 0$.



Brute Force Algorithm

Algorithm 1: Weighted Interval Scheduling – Brute Force

Input: $n; s_1, \dots, s_n; f_1, \dots, f_n; w_1, \dots, w_n;$

Output: Optimal weight $OPT(n)$.

- 1 Sort jobs by finish times so that $f_1 \leq f_2 \leq \dots \leq f_n$;
- 2 Compute $p(1), p(2), \dots, p(n)$;
- 3 **return** B-Sched (n);

Algorithm 2: B-Sched (j)

- 1 **if** $j = 0$ **then**
- 2 **return** 0;
- 3 **else**
- 4 **return** $\max\{w_j + B-Sched(p(j)), B-Sched(j - 1)\}$;

Binary Choice

Recurrence template: $OPT(j) = \text{value of optimal solution to the problem consisting of job requests } 1, 2, \dots, j.$

Optimal substructure:

Case 1: OPT selects job j .

- o collect weight w_j ,
- o can't use incompatible jobs $\{p(j) + 1, p(j) + 2, \dots, j - 1\}$,
- o must include optimal solution to problem consisting of remaining compatible jobs $1, 2, \dots, p(j)$.

Case 2: OPT does not select job j .

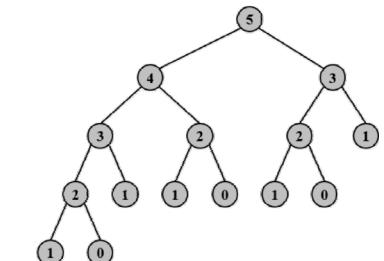
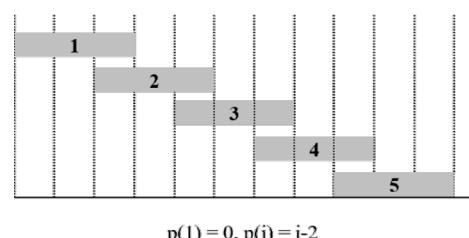
- o must include optimal solution to problem consisting of remaining compatible jobs $1, 2, \dots, j - 1$.

$$OPT(j) = \begin{cases} 0, & j = 0, \\ \max\{w_j + OPT(p(j)), OPT(j - 1)\}, & \text{otherwise} \end{cases}$$

Brute Force Algorithm

Observation: Recursive algorithm fails spectacularly because of redundant sub-problems \Rightarrow **exponential algorithms**.

Example: Number of recursive calls for family of "layered" instances grows like Fibonacci sequence.



Memoization: Store sub-results in cache; lookup as needed

Algorithm 3: Weighted Interval Scheduling – Memoization

Input: $n; s_1, \dots, s_n; f_1, \dots, f_n; w_1, \dots, w_n;$

Output: Optimal weight $OPT(n)$.

```

1 Sort jobs by finish times so that  $f_1 \leq f_2 \leq \dots \leq f_n$ ;
2 Compute  $p(1), p(2), \dots, p(n)$ ;
3  $M[0] = 0$ ; // global array
4 return M-Sched ( $n$ );

```

Algorithm 4: M-Sched (j)

```

1 if  $M[j]$  is uninitialized then
2    $\lfloor M[j] = \max\{w_j + M\text{-Sched}(p(j)), M\text{-Sched}(j - 1)\}$ ;
3 return  $M[j]$ ;

```

Finding a Solution from the OPT Value

Algorithm 5: Find-Solution (j)

```

1 if  $j = 0$  then
2   return  $\emptyset$ ;
3 else if  $w_j + M[p(j)] > M[j - 1]$  then
4   return  $\{j\} \cup \text{Find-Solution}(p(j))$ ;
5 else
6    $\lfloor \text{Return Find-Solution}(j - 1)$ ;

```

- Run $\text{Find-Solution}(n)$ to find optimal schedule;
- # of recursive calls $1 \leq n \Rightarrow O(n)$;

Running Time

Claim: Memoized version of algorithm takes $O(n \log n)$ time.

- Sort by finish time: $O(n \log n)$.
- Computing $p(\cdot)$: $O(n \log n)$ via sorting by start time.
- M-Sched(j): each invocation takes $O(1)$ time and either
 - (1) returns an existing value $M[j]$
 - (2) initializes $M[j]$ and makes two recursive calls
- Progress measure $\Phi = \text{number nonempty entries of } M[\cdot]$.
 - ▷ initially $\Phi = 0$, throughout $\Phi \leq n$.
 - ▷ (2) increases Φ by 1 \Rightarrow at most $2n$ recursive calls.
- Overall running time of M-Sched(n) is $O(n)$.

Remark: $O(n)$ if jobs are pre-sorted by start and finish times.

Tabulation: Bottom-Up Dynamic Programming

Algorithm 6: Weighted Interval Scheduling – Tabulation

Input: $n; s_1, \dots, s_n; f_1, \dots, f_n; w_1, \dots, w_n;$

Output: Optimal weight $OPT(n)$.

```

1 Sort jobs by finish times so that  $f_1 \leq f_2 \leq \dots \leq f_n$ ;
2 Compute  $p(1), p(2), \dots, p(n)$ ;
3  $M[0] = 0$ ;
4 for  $j = 1 \rightarrow n$  do
5    $\lfloor M[j] = \max\{w_j + M[p(j)], M[j - 1]\}$ ;

```

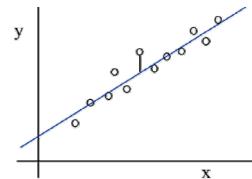
Running Time: $O(n \log n)$.

Those who cannot remember the past are condemned to repeat it.

— Kevin Wayne@Princeton

Segmented Least Squares

- Foundational problem in statistic and numerical analysis.
- Given n points in the plane: $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$.
- Find a line $y = ax + b$ to minimize the sum of the squared error:

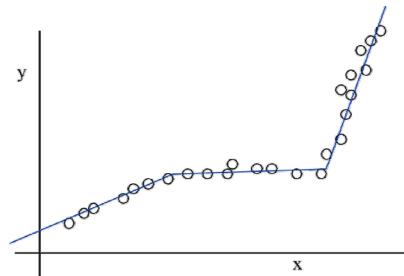


Solution: Calculus \Rightarrow min error is achieved when

$$a = \frac{n \sum_i x_i y_i - (\sum_i x_i)(\sum_i y_i)}{n \sum_i x_i^2 - (\sum_i x_i)^2}, b = \frac{\sum_i y_i - a \sum_i x_i}{n}$$

Segmented Least Squares

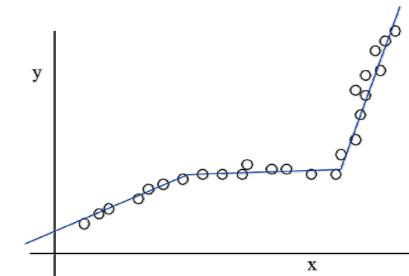
- Points lie roughly on a sequence of several line segments.
- Given n points in the plane: $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$ with $x_1 < x_2 < \dots < x_n$, find a sequence of lines that minimizes:
 - the sum of the sums of the squared errors E in each segment
 - the number of lines L
- Tradeoff function: $E + cL$, for some constant $c > 0$.



Segmented Least Squares

- Points lie roughly on a sequence of several line segments.
- Given n points in the plane: $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$ with $x_1 < x_2 < \dots < x_n$, find a sequence of lines that minimizes $f(x)$.

Question: What's a reasonable choice for $f(x)$ to balance accuracy (goodness of fit) and parsimony (number of lines)?



Multiway Choice

Notation:

- $OPT(j)$ = minimum cost for points p_1, p_{i+1}, \dots, p_j .
- $e(i, j)$ = minimum sum of squares for points p_i, p_{i+1}, \dots, p_j .

Compute $OPT(j)$:

- Last segment uses points p_i, p_{i+1}, \dots, p_j for some i .
- $Cost = e(i, j) + c + OPT(i - 1)$.

$$OPT(j) = \begin{cases} 0, & j = 0, \\ \min_{1 \leq i \leq j} \{e(i, j) + c + OPT(i - 1)\}, & \text{otherwise} \end{cases}$$

Segmented Least Squares

Algorithm 7: Segmented Square Error (SSE)

Input: $n; p_1, \dots, p_n; c;$

Output: Optimal square error for $p_1, \dots, p_n.$

```

1 for  $j = 1 \rightarrow n$  do
2   for  $i = 1 \rightarrow j$  do
3     compute least square error  $e_{ij}$  for segment  $p_i, \dots, p_j;$ 
4    $M[0] = 0;$ 
5   for  $j = 1 \rightarrow n$  do
6      $M[j] = \min_{1 \leq i \leq j} \{e_{ij} + c + M[i - 1]\};$ 
7 return  $M[n];$ 
```

Time Complexity: $O(n^3)$ (can be improved to $O(n^2)$)

Space Complexity: $O(n^2).$

Knapsack Problem

Given n objects and a "knapsack".

Item i weighs $w_i > 0$ kilograms and has value $v_i > 0.$

Knapsack has capacity of W kilograms.

Goal: fill knapsack so as to maximize total value.

Example: $\{3, 4\}$ has value 40.

$W = 11$

#	value	weight
1	1	1
2	6	2
3	18	5
4	22	6
5	28	7

Greedy: repeatedly add item with maximum ratio $v_i/w_i.$

Example: $\{5, 2, 1\}$ achieves only value = 35 \Rightarrow greedy not optimal.

Algorithm Analysis

Theorem (Bellman, 1961) SSE solves the segmented least squares problem in $O(n^3)$ time and $O(n^2)$ space.

Proof: Bottleneck = computing e_{ij} for $O(n^2)$ pairs,

$$a = \frac{n \sum_i x_i y_i - (\sum_i x_i)(\sum_i y_i)}{n \sum_i x_i^2 - (\sum_i x_i)^2}, \quad b = \frac{\sum_i y_i - a \sum_i x_i}{n},$$

$O(n)$ per pair e_{ij} using previous formula. \square

Remark: Can be improved to $O(n^2)$ time.

- o $\forall i$: precompute cumulative sums $\sum_{k=1}^i x_k, \sum_{k=1}^i y_k, \sum_{k=1}^i x_k^2, \sum_{k=1}^i x_k y_k,$
- o Using cumulative sums, we can compute e_{ij} in $O(1)$ time.

First Attempt

Definition: $OPT(i) = \max$ profit subset of items $1, \dots, i.$

Case 1: OPT does not select item $i.$

- o OPT selects best of $\{1, 2, \dots, i - 1\}.$

Case 2: OPT selects item $i.$

- o accepting item i does not immediately imply that we will have to reject other items,
- o without knowing what other items were selected before i , we don't even know if we have enough room for $i.$

Conclusion: Need more sub-problems!

Adding a New Variable

Definiton: $OPT(i)$ = max profit subset of items $1, \dots, i$ with weight limit w .

Case 1: OPT does not select item i .

- OPT selects best of $\{1, 2, \dots, i-1\}$ using weight limit w

Case 2: OPT selects item i .

- new weight limit = $w - w_i$
- OPT selects best of using $\{1, 2, \dots, i-1\}$ this new weight limit

$$OPT(i, w) = \begin{cases} 0, & j = 0, \\ OPT(i-1, w), & w_i > w, \\ \max\{OPT(i-1, w), v_i + OPT(i-1, w - w_i)\}, & \text{otherwise} \end{cases}$$

Knapsack Algorithm

	W + 1											
	0	1	2	3	4	5	6	7	8	9	10	11
ϕ	0	0	0	0	0	0	0	0	0	0	0	0
{ 1 }	0	1	1	1	1	1	1	1	1	1	1	1
{ 1, 2 }	0	1	6	7	7	7	7	7	7	7	7	7
{ 1, 2, 3 }	0	1	6	7	7	18	19	24	25	25	25	25
{ 1, 2, 3, 4 }	0	1	6	7	7	18	22	24	28	29	29	40
{ 1, 2, 3, 4, 5 }	0	1	6	7	7	18	22	28	29	34	34	40

OPT: {4, 3}

value = $22 + 18 = 40$

Item	Value	Weight
1	1	1
2	6	2
3	18	5
4	22	6
5	28	7

W = 11

Bottom-Up Algorithm (Fill up an n -by- W array)

Algorithm 8: Knapsack Algorithm using n -by- W Array

Input: $n, W, w_1, \dots, w_n, v_1, \dots, v_n$;

Output: Optimal value of knapsack with W .

```

1 for  $w = 0 \rightarrow W$  do
2    $M[0, w] = 0$ ;
3 for  $i = 1 \rightarrow n$  do
4   for  $w = 1 \rightarrow W$  do
5     if  $w_i > w$  then
6        $M[i, w] = M[i - 1, w]$ ;
7     else
8        $M[i, w] = \max\{M[i - 1, w], v_i + M[i - 1, w - w_i]\}$ ;
9 return  $M[n, W]$ ;
```

Running Time

Running time: $\Theta(nW)$.

- Not polynomial in input size!
- "Pseudo-polynomial".
- Decision version of Knapsack is NP-complete.

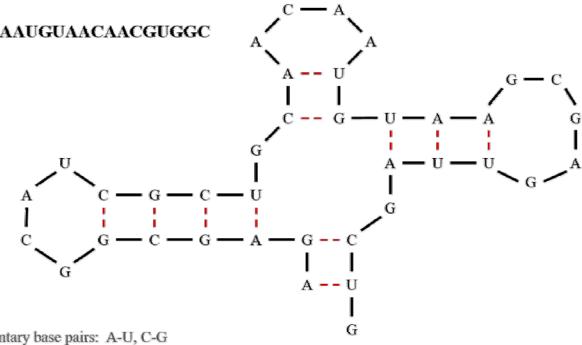
Knapsack approximation algorithm: There exists a poly-time algorithm that produces a feasible solution that has value within 0.01% of optimum.

RNA Secondary Structure

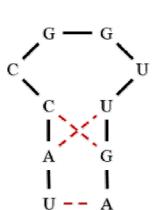
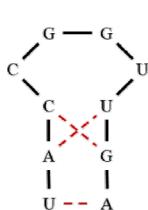
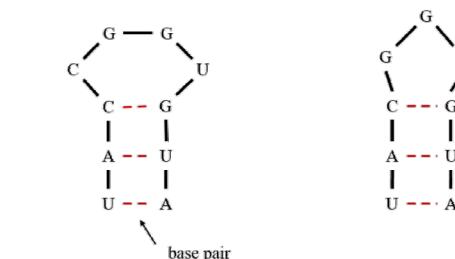
RNA: String $B = b_1 b_2 \cdots b_n$ over alphabet $\{A, C, G, U\}$.

Secondary structure: RNA is single-stranded so it tends to loop back and form base pairs with itself. This structure is essential for understanding behavior of molecule.

Example: GUCAUUGAGCGAAUGUAACACGUGGC
UACGGCGAGA



Examples



ok



sharp turn



crossing

RNA Secondary Structure

Secondary structure: A set of pairs $S = \{(b_i, b_j)\}$ that satisfy:

[Watson-Crick] S is a matching and each pair in S is a Watson-Crick complement: A-U, U-A, C-G, or G-C.

[No sharp turns] The ends of each pair are separated by at least 4 intervening bases. If $(b_i, b_j) \in S$, then $i < j - 4$.

[Non-crossing] If (b_i, b_j) and (b_k, b_l) are two pairs in S , then we cannot have $i < k < j < l$.

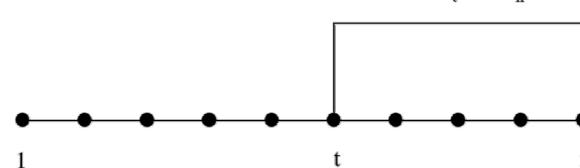
Free energy: Usual hypothesis is that an RNA molecule will form the secondary structure with the optimum total free energy.

Goal: Given an RNA molecule $B = b_1 b_2 \cdots b_n$, find a secondary structure S that maximizes the number of base pairs

Subproblems

First attempt: $OPT(j) = \text{maximum number of base pairs in a secondary structure of the substring } b_1 b_2 \cdots b_j$.

match b_t and b_n



Difficulty: Results in two sub-problems.

- Finding secondary structure in: $b_1 b_2 \cdots b_{t-1}$.
- Finding secondary structure in: $b_{t+1} b_{t+2} \cdots b_{n-1}$.

Dynamic Programming Over Intervals

Notation: $OPT(j)$ = maximum number of base pairs in a secondary structure of the substring $b_i b_{i+1} \dots b_j$.

Case 1: If $i \geq j - 4$.

- $OPT(i, j) = 0$ by no-sharp turns condition.

Case 2: Base b_j is not involved in a pair.

- $OPT(i, j) = OPT(i, j - 1)$

Case 3: Base b_j pairs with b_t for some $i \leq t < j - 4$.

- non-crossing constraint decouples resulting sub-problems
- $OPT(i, j) = 1 + \max_t \{OPT(i, t - 1) + OPT(t + 1, j - 1)\}$

Remark: Same core idea in CKY algorithm to parse context-free grammars.

Dynamic Programming Summary

Recipe

- Characterize structure of problem.
- Recursively define value of optimal solution.
- Compute value of optimal solution.
- Construct optimal solution from computed information.

Dynamic programming techniques

- Binary choice: weighted interval scheduling.
- Multi-way choice: segmented least squares.
- Adding a new variable: knapsack.
- Dynamic programming over interval

Top-down vs. bottom-up: different people have different intuitions.

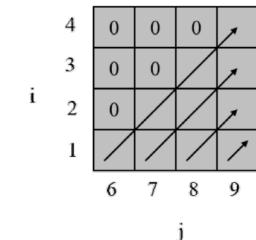
Bottom Up Dynamic Programming Over Intervals

Question: What order to solve the sub-problems?

Answer: Do shortest intervals first.

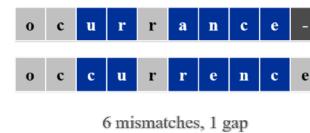
```
RNA(b1,...,bn) {
    for k = 5, 6, ..., n-1
        for i = 1, 2, ..., n-k
            j = i + k
            Compute M[i, j]
    return M[1, n]
```

using recurrence

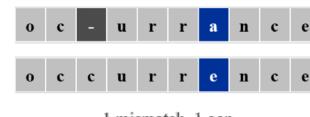


Running time: $O(n^3)$.

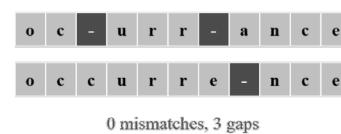
String Similarity: How similar are two strings?



6 mismatches, 1 gap



1 mismatch, 1 gap



0 mismatches, 3 gaps

How similar are two strings?

- occurrence
- occurrence

Linear Space

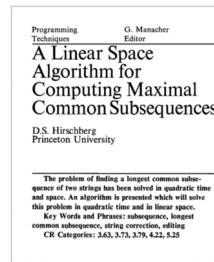
Question: Can we avoid using quadratic space?

Easy. Optimal value in $O(m + n)$ space* and $O(mn)$ time.

- Compute $OPT(i, \cdot)$ from $OPT(i - 1, \cdot)$.
- No longer a simple way to recover alignment itself.

Theorem. [Hirschberg 1975] Optimal alignment in $O(m + n)$ space and $O(mn)$ time.

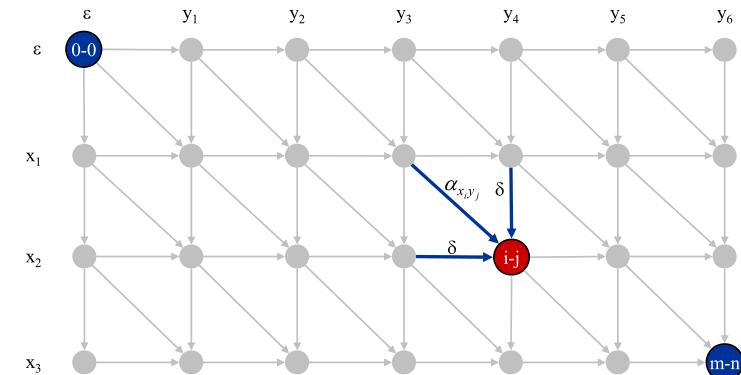
- Clever combination of divide-and-conquer and dynamic programming.
- Inspired by idea of Savitch from complexity theory.



*including space storing original strings

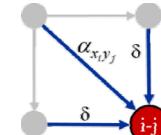
Edit Distance Graph

- Let $f(i, j)$ be shortest path from $(0, 0)$ to (i, j) .
- Observation: $f(i, j) = OPT(i, j)$.



Edit Distance Graph

- Let $f(i, j)$ be shortest path from $(0, 0)$ to (i, j) .
- Observation: $f(i, j) = OPT(i, j)$.



Proof: (by strong induction on $i + j$)

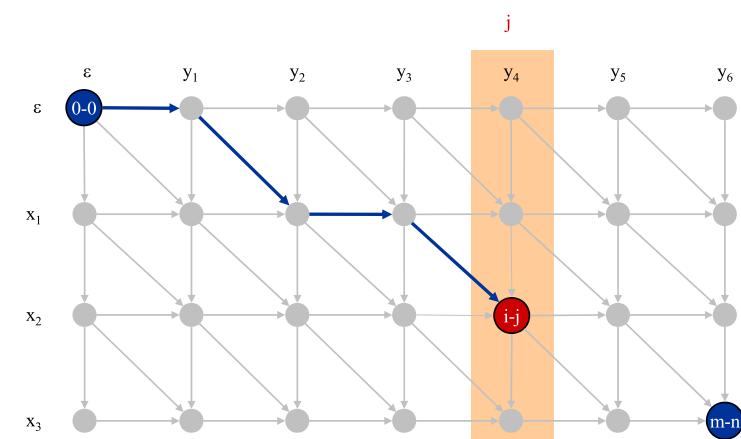
Base case: $f(0, 0) = OPT(0, 0) = 0$

Inductive hypothesis: assume true for all (i', j') with $i' + j' < i + j$.

Induction: Last edge on shortest path to (i, j) is from $(i - 1, j - 1)$, $(i - 1, j)$, or $(i, j - 1)$.

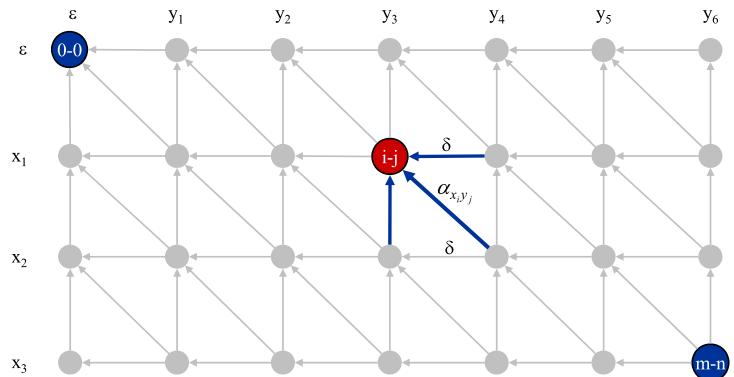
$$\begin{aligned} f(i, j) &= \min\{a_{x_i, y_j} + f(i - 1, j - 1), \delta + f(i - 1, j), \delta + f(i, j - 1)\} \\ &= \min\{a_{x_i, y_j} + OPT(i - 1, j - 1), \delta + OPT(i - 1, j), \delta + OPT(i, j - 1)\} \\ &= OPT(i, j) \end{aligned}$$

- Let $f(i, j)$ be shortest path from $(0, 0)$ to (i, j) .
- Can compute $f(\cdot, j)$ for any j in $O(mn)$ time and $O(m + n)$ space.



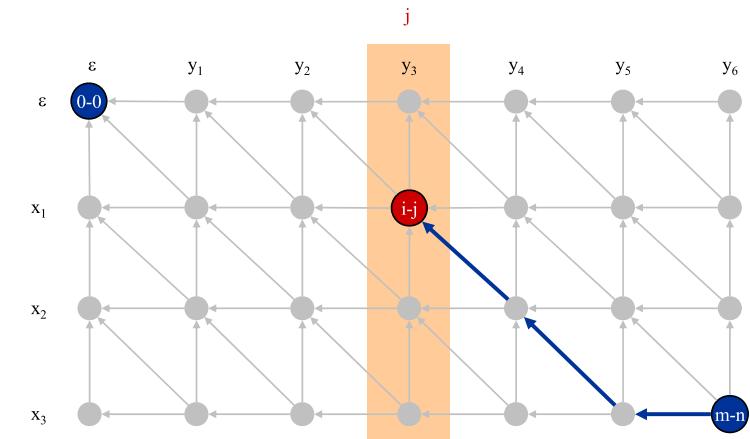
Edit Distance Graph

- Let $g(i,j)$ be shortest path from (i,j) to (m,n) .
- Can compute by reversing the edge orientations and inverting the roles of $(0,0)$ and (m,n)



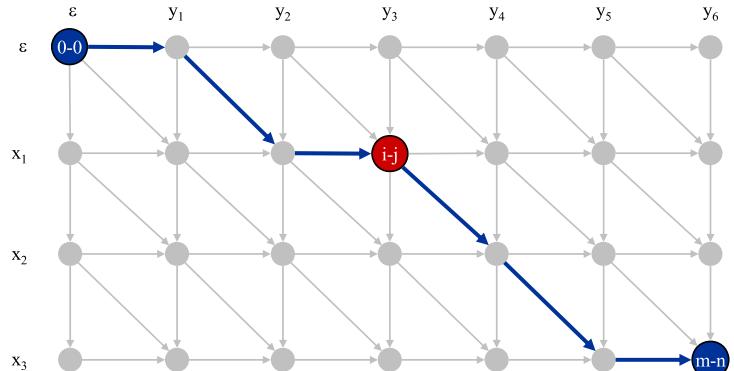
Edit Distance Graph

- Let $g(i,j)$ be shortest path from (i,j) to (m,n) .
- Can compute $g(\cdot, j)$ for any j in $O(mn)$ time and $O(m + n)$ space.



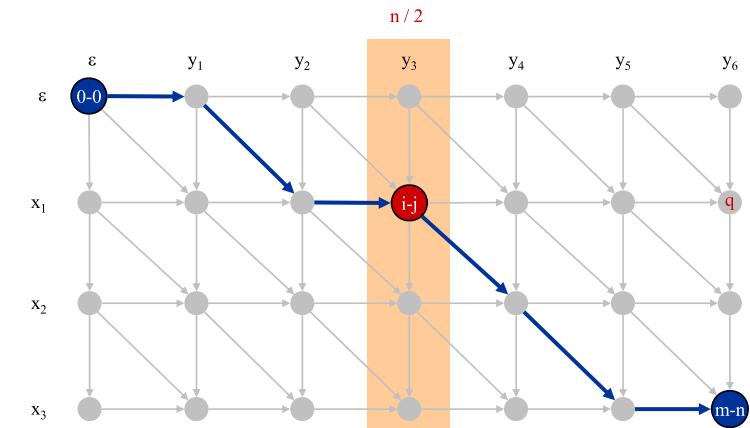
Edit Distance Graph

Observation 1: The cost of the shortest path that uses (i,j) is $f(i,j) + g(i,j)$.



Edit Distance Graph

Observation 2: Let q be an index that minimizes $f(q, n/2) + g(q, n/2)$. Then, the shortest path from $(0, 0)$ to (m, n) uses $(q, n/2)$.

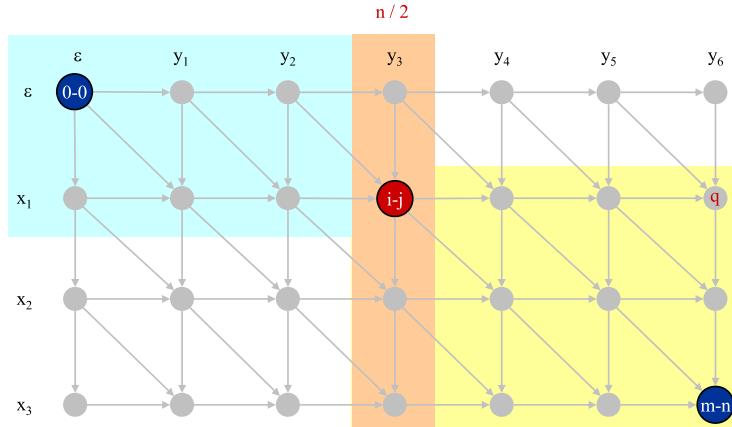


Edit Distance Graph

Divide: find index q that minimizes $f(q, n/2) + g(q, n/2)$ using DP.

Align x_q and $y_{n/2}$.

Conquer: recursively compute optimal alignment in each piece.



Running Time Analysis

Theorem. Let $T(m, n) = \max$ running time of algorithm on strings of length m and n . $T(m, n) = O(mn)$

Proof: (by induction on n)

- $O(mn)$ time to compute $f(\cdot, n/2)$ and $g(\cdot, n/2)$ and find index q .
- $T(q, n/2) + T(m - q, n/2)$ time for two recursive calls
- Choose constant c so that:

$$T(m, 2) \leq cm$$

$$T(2, n) \leq cn$$

$$T(m, n) \leq cmn + T(q, n/2) + T(m - q, n/2)$$

Running Time Analysis Warmup

Theorem: Let $T(m, n) = \max$ running time of algorithm on strings of length at most m and n . $T(m, n) = O(mn \log n)$.

$$T(m, n) \leq 2T(m, n/2) + O(mn) \Rightarrow T(m, n) = O(mn \log n)$$

Remark: Analysis is not tight because two sub-problems are of size $(q, n/2)$ and $(m - q, n/2)$. In next slide, we save $\log n$ factor.

Running Time Analysis (Continued)

Theorem. Let $T(m, n) = \max$ running time of algorithm on strings of length m and n . $T(m, n) = O(mn)$

Proof:

- Base cases: $m = 2$ or $n = 2$.
- Inductive hypothesis: $T(m, n) \leq 2cmn$.

$$\begin{aligned} T(m, n) &\leq T(q, n/2) + T(m - q, n/2) + cmn \\ &\leq 2cqn/2 + 2c(m - q)n/2 + cmn \\ &= cqn + cmn - cqn + cmn \\ &= 2cmn \end{aligned}$$